

Augmented gradient descent (LR = 0.200, σ =[None, 3, 7, 11])
prompt = "a coffee mug on the floor of an office"
(CLIP version = ViT-L/14)

